

目 录

1	分析 resize 流程前的必要知识	2
1.1	nova 中的 RPC 机制	2
1.2	重要的数据类型	2
1.2.1	req	2
1.2.2	context	2
1.2.3	instance	2
2	nova-api 阶段	2
3	nova-conductor 部分	3
4	冷迁移	4
4.1	冷迁移中的 nova-conductor 部分	4
4.2	冷迁移中的 nova-compute 部分	5
4.2.1	目的主机上的操作: prep_resize	5
4.2.2	源主机的操作: resize_instance	5
4.2.3	目的主机上的操作: finish_resize	6
5	在线迁移	7
5.1	在线迁移中的 nova-conductor 部分	7
5.2	在线迁移中的 nova-compute 部分	8
5.2.1	目的主机上的操作: pre_live_migration	9
5.2.2	源主机上的操作: live_migration	12

1 分析 resize 流程前的必要知识

1.1 nova 中的 RPC 机制

1.2 重要的数据类型

1.2.1 req

1.2.2 context

```
1 # 根据req创建环境上下文context
2 # context是nova/context.py中的RequestContext类
3 context = req.environ["nova.context"]
```

1.2.3 instance

```
1 # 根据req和instance_id创建instance
2 # instance是nova/context/instance.py中的Instance类
3 instance = self._get_server(context, req, instance_id)
4
5
6 instance_type
7 flavor_id
8 deltas
9 quotas
10 vm_state
```

2 nova-api 阶段

入口函数为：

```
1 # 这个函数在nova/api/openstack/compute/servers.py
2 def _resize(self, req, instance_id, flavor_id, **kwargs):
3     ...
4     try:
5         # compute_api是nova/compute/api.py中的API类
6         self.compute_api.resize(context, instance, flavor_id, **kwargs)
7     ...
```

进一步看 API.resize() 函数：

```
1 # nova/compute/api.py API.resize()
2 def resize(self, context, instance, flavor_id=None,
3            **extra_instance_updates):
```

```

4      ...
5
6      # filter_properties与选择本地扩容或选择异地扩容有关
7      filter_properties = {'ignore_hosts': []}
8
9      if not CONF.allow_resize_to_same_host:
10         filter_properties['ignore_hosts'].append(instance['host'])
11
12      if (not flavor_id and not CONF.allow_migrate_to_same_host):
13         filter_properties['ignore_hosts'].append(instance['host'])
14
15      ...
16
17      # scheduler_hint挺重要的, 是nova-scheduler的参数
18      scheduler_hint = {'filter_properties': filter_properties}
19      self.compute_task_api.resize_instance(context, instance,
20         extra_instance_updates, scheduler_hint=scheduler_hint,
21         flavor=new_instance_type,
22         reservations=quotas.reservations or [])

```

```

1      # nova/conductor/api.py ComputeTaskAPI.resize_instance()
2      def resize_instance(self, context, instance, extra_instance_updates,
3         scheduler_hint, flavor, reservations):
4         self.conductor_compute_rpcapi.migrate_server(
5             context, instance, scheduler_hint, False, False, flavor,
6             None, None, reservations)

```

```

1      # nova/conductor/rpcapi.py ComputeTaskAPI.migrate_server()
2      def migrate_server(self, context, instance, scheduler_hint, live, rebuild,
3         flavor, block_migration, disk_over_commit,
4         reservations=None):
5         ...
6         cctxt = self.client.prepare(version=version)
7         return cctxt.call(context, 'migrate_server',
8             instance=instance, scheduler_hint=scheduler_hint,
9             live=live, rebuild=rebuild, flavor=flavor_p,
10             block_migration=block_migration,
11             disk_over_commit=disk_over_commit,
12             reservations=reservations)

```

3 nova-conductor 部分

```

1      # nova/conductor/manager.py ComputeTaskManager.migrate_server()
2      def migrate_server(self, context, instance, scheduler_hint, live, rebuild,
3         flavor, block_migration, disk_over_commit, reservations=None):
4         ...
5         if live and not rebuild and not flavor:
6             self._live_migrate(context, instance, scheduler_hint,
7                 block_migration, disk_over_commit)
8         elif not live and not rebuild and flavor:

```

```

9         ...
10         with compute_utils.EventReporter(context, 'cold_migrate',
11                                           instance_uuid):
12             self._cold_migrate(context, instance, flavor,
13                               scheduler_hint['filter_properties'],
14                               reservations)
15         ...

```

4 冷迁移

4.1 冷迁移中的 nova-conductor 部分

```

1  # nova/conductor/manager.py ComputeTaskManager._cold_migrate()
2  def _cold_migrate(self, context, instance, flavor, filter_properties,
3                    reservations):
4      ...
5      try:
6          ...
7          # 选择目的主机
8          hosts = self.scheduler_client.select_destinations(
9                  context, request_spec, filter_properties)
10         host_state = hosts[0]
11         ...
12
13     try:
14         ...
15         (host, node) = (host_state['host'], host_state['nodename'])
16         self.compute_rpcapi.prep_resize(
17             context, image, instance,
18             flavor, host,
19             reservations, request_spec=request_spec,
20             filter_properties=filter_properties, node=node)
21     ...

```

```

1  # nova/compute/rpcapi.py ComputeAPI.prep_resize()
2  def prep_resize(self, ctxt, image, instance, instance_type, host,
3                  reservations=None, request_spec=None,
4                  filter_properties=None, node=None):
5      ...
6      cctxt = self.client.prepare(server=host, version=version)
7      cctxt.cast(ctxt, 'prep_resize',
8                 instance=instance,
9                 instance_type=instance_type_p,
10                 image=image_p, reservations=reservations,
11                 request_spec=request_spec,
12                 filter_properties=filter_properties,
13                 node=node)

```

4.2 冷迁移中的 nova-compute 部分

4.2.1 目的主机上的操作: `prep_resize`

```

1  # nova/compute/manager.py ComputeManager.prep_resize()
2  def prep_resize(self, context, image, instance, instance_type,
3                  reservations, request_spec, filter_properties, node):
4      ...
5      with self._error_out_instance_on_exception(context, instance,
6                                                  quotas=quotas):
7          ...
8          try:
9              self._prep_resize(context, image, instance,
10                              instance_type, quotas,
11                              request_spec, filter_properties,
12                              node)
13      ...

```

```

1  # nova/compute/manager.py ComputeManager._prep_resize()
2  def _prep_resize(self, context, image, instance, instance_type,
3                  quotas, request_spec, filter_properties, node):
4      ...
5      with rt.resize_claim(context, instance, instance_type,
6                           image_meta=image, limits=limits) as claim:
7          ...
8          self.compute_rpcapi.resize_instance(
9              context, instance, claim.migration, image,
10              instance_type, quotas.reservations)
11

```

```

1  # nova/compute/rpcapi.py ComputeAPI.resize_instance()
2  def resize_instance(self, ctxt, instance, migration, image, instance_type,
3                     reservations=None):
4      ...
5      cctxt = self.client.prepare(server=_compute_host(None, instance),
6                                  version=version)
7      cctxt.cast(ctxt, 'resize_instance',
8                 instance=instance, migration=migration,
9                 image=image, reservations=reservations,
10                 instance_type=instance_type_p)

```

4.2.2 源主机的操作: `resize_instance`

```

1  # nova/compute/manager.py ComputeManager.resize_instance()
2  def resize_instance(self, context, instance, image,
3                    reservations, migration, instance_type,
4                    clean_shutdown=True):
5      ...
6      with self._error_out_instance_on_exception(context, instance,
7                                                  quotas=quotas):

```

```

8      ...
9      # 获得虚拟机块设备的信息
10     block_device_info = self._get_instance_block_device_info(
11         context, instance, bdms=bdms)
12
13     # 关闭虚拟机并迁移虚拟机的增量文件
14     disk_info = self.driver.migrate_disk_and_power_off(
15         context, instance, migration.dest_host,
16         instance_type, network_info,
17         block_device_info,
18         timeout, retry_interval)
19     ...
20     self.compute_rpcapi.finish_resize(context, instance,
21         migration, image, disk_info,
22         migration.dest_compute, reservations=quotas.reservations)
23     ...

```

migrate_disk_and_power_off() 是源主机上将虚拟机迁移给目的主机的实现函数，主要利用了 libvirt API。这个函数的分析在《nova 调用 libvirt》中的“nova 扩容时对 libvirt 的调用”一节。

```

1  # nova/compute/rpcapi.py(690) ComputeAPI.finish_resize()
2  def finish_resize(self, ctxt, instance, migration, image, disk_info,
3      host, reservations=None):
4      ...
5      cctxt = self.client.prepare(server=host, version=version)
6      cctxt.cast(ctxt, 'finish_resize',
7          instance=instance, migration=migration,
8          image=image, disk_info=disk_info, reservations=reservations)

```

4.2.3 目的主机上的操作：finish_resize

```

1  # nova/compute/manager.py ComputeManager.finish_resize()
2  def finish_resize(self, context, disk_info, image, instance,
3      reservations, migration):
4      quotas = quotas_obj.Quotas.from_reservations(context,
5          reservations,
6          instance=instance)
7      try:
8          self._finish_resize(context, instance, migration,
9              disk_info, image)
10     ...

```

```

1  # nova/compute/manager.py ComputeManager._finish_resize()
2  def _finish_resize(self, context, instance, migration, disk_info,
3      image):
4      ...
5      try:
6          self.driver.finish_migration(context, migration, instance,
7              disk_info,
8              network_info,

```

```

9         image, resize_instance,
10         block_device_info, power_on)
11     ...

```

5 在线迁移

5.1 在线迁移中的 nova-conductor 部分

```

1  # nova/conductor/manager.py ComputeTaskManager._live_migrate()
2  def _live_migrate(self, context, instance, scheduler_hint,
3                  block_migration, disk_over_commit):
4      destination = scheduler_hint.get("host")
5      try:
6          live_migrate.execute(context, instance, destination,
7                              block_migration, disk_over_commit)
8      ...

```

```

1  # nova/conductor/tasks/live_migrate.py execute()
2  def execute(context, instance, destination,
3             block_migration, disk_over_commit):
4      task = LiveMigrationTask(context, instance,
5                              destination,
6                              block_migration,
7                              disk_over_commit)
8  # TODO(johngarbutt) create a superclass that contains a safe_execute call
9  return task.execute()

```

```

1  # nova/conductor/tasks/live_migrate.py LiveMigrationTask.execute()
2  def execute(self):
3      self._check_instance_is_running()
4      self._check_host_is_up(self.source)
5
6      if not self.destination:
7          self.destination = self._find_destination()
8      else:
9          self._check_requested_destination()
10
11     # TODO(johngarbutt) need to move complexity out of compute manager
12     # TODO(johngarbutt) disk_over_commit?
13     return self.compute_rpcapi.live_migration(self.context,
14        host=self.source,
15        instance=self.instance,
16        dest=self.destination,
17        block_migration=self.block_migration,
18        migrate_data=self.migrate_data)

```

```

1  # nova/compute/rpcapi.py ComputeAPI.live_migration()
2  def live_migration(self, ctxt, instance, dest, block_migration, host,

```

```

3         migrate_data=None):
4         ...
5         cctxt = self.client.prepare(server=host, version=version)
6         cctxt.cast(ctxt, 'live_migration', instance=instance,
7                     dest=dest, block_migration=block_migration,
8                     migrate_data=migrate_data)

```

5.2 在线迁移中的 nova-compute 部分

```

1  # nova/compute/manager.py ComputeManager.live_migration()
2  def live_migration(self, context, dest, instance, block_migration,
3                    migrate_data):
4      """Executing live migration.
5
6      :param context: security context
7      :param instance: a nova.objects.instance.Instance object
8      :param dest: destination host
9      :param block_migration: if true, prepare for block migration
10     :param migrate_data: implementation specific params
11
12     """
13
14     # NOTE(danms): since instance is not the first parameter, we can't
15     # use @object_compat on this method. Since this is the only example,
16     # we do this manually instead of complicating the decorator
17     if not isinstance(instance, obj_base.NovaObject):
18         expected = ['metadata', 'system_metadata',
19                     'security_groups', 'info_cache']
20         instance = objects.Instance._from_db_object(
21             context, objects.Instance(), instance,
22             expected_attrs=expected)
23
24     # Create a local copy since we'll be modifying the dictionary
25     migrate_data = dict(migrate_data or {})
26     try:
27         if block_migration:
28             block_device_info = self._get_instance_block_device_info(
29                 context, instance)
30             disk = self.driver.get_instance_disk_info(
31                 instance.name, block_device_info=block_device_info)
32         else:
33             disk = None
34
35         pre_migration_data = self.compute_rpcapi.pre_live_migration(
36             context, instance,
37             block_migration, disk, dest, migrate_data)
38         migrate_data['pre_live_migration_result'] = pre_migration_data
39
40     except Exception:
41         with excutils.save_and_reraise_exception():
42             LOG.exception(_LE('Pre live migration failed at %s'),
43                           dest, instance=instance)
44             self._rollback_live_migration(context, instance, dest,

```



```

45         block_migration, migrate_data)
46
47     # Executing live migration
48     # live_migration might raises exceptions, but
49     # nothing must be recovered in this version.
50     self.driver.live_migration(context, instance, dest,
51                               self._post_live_migration,
52                               self._rollback_live_migration,
53                               block_migration, migrate_data)

```

5.2.1 目的主机上的操作: pre_live_migration

```

1  def pre_live_migration(self, context, instance, block_migration, disk,
2                        migrate_data):
3      """Preparations for live migration at dest host.
4
5      :param context: security context
6      :param instance: dict of instance data
7      :param block_migration: if true, prepare for block migration
8      :param migrate_data: if not None, it is a dict which holds data
9                          required for live migration without shared
10                         storage.
11
12      """
13     block_device_info = self._get_instance_block_device_info(
14         context, instance, refresh_conn_info=True)
15
16     network_info = self._get_instance_nw_info(context, instance)
17     self._notify_about_instance_usage(
18         context, instance, "live_migration.pre.start",
19         network_info=network_info)
20
21     pre_live_migration_data = self.driver.pre_live_migration(context,
22                                                             instance,
23                                                             block_device_info,
24                                                             network_info,
25                                                             disk,
26                                                             migrate_data)
27
28     # NOTE(tr3buchet): setup networks on destination host
29     self.network_api.setup_networks_on_host(context, instance,
30                                             self.host)
31
32     # Creating filters to hypervisors and firewalls.
33     # An example is that nova-instance-instance-xxx,
34     # which is written to libvirt.xml(Check "virsh nwfilter-list")
35     # This nwfilter is necessary on the destination host.
36     # In addition, this method is creating filtering rule
37     # onto destination host.
38     self.driver.ensure_filtering_rules_for_instance(instance,
39                                                     network_info)
40
41     self._notify_about_instance_usage(

```

```

42         context, instance, "live_migration.pre.end",
43         network_info=network_info)
44
45     return pre_live_migration_data

```

```

1     # nova/virt/libvirt/driver.py LibvirtDriver.pre_live_migration()
2     def pre_live_migration(self, context, instance, block_device_info,
3         network_info, disk_info, migrate_data=None):
4         """Preparation live migration."""
5         # Steps for volume backed instance live migration w/o shared storage.
6         is_shared_block_storage = True
7         is_shared_instance_path = True
8         is_block_migration = True
9         instance_relative_path = None
10        if migrate_data:
11            is_shared_block_storage = migrate_data.get(
12                'is_shared_block_storage', True)
13            is_shared_instance_path = migrate_data.get(
14                'is_shared_instance_path', True)
15            is_block_migration = migrate_data.get('block_migration', True)
16            instance_relative_path = migrate_data.get('instance_relative_path')
17
18        if not (is_shared_instance_path and is_shared_block_storage):
19            # NOTE(dims): Using config drive with iso format does not work
20            # because of a bug in libvirt with read only devices. However
21            # one can use vfat as config_drive_format which works fine.
22            # Please see bug/1246201 for details on the libvirt bug.
23            if CONF.config_drive_format != 'vfat':
24                if configdrive.required_by(instance):
25                    raise exception.NoLiveMigrationForConfigDriveInLibVirt()
26
27        if not is_shared_instance_path:
28            # NOTE(mikal): this doesn't use libvirt_utils.get_instance_path
29            # because we are ensuring that the same instance directory name
30            # is used as was at the source
31            if instance_relative_path:
32                instance_dir = os.path.join(CONF.instances_path,
33                                            instance_relative_path)
34            else:
35                instance_dir = libvirt_utils.get_instance_path(instance)
36
37            if os.path.exists(instance_dir):
38                raise exception.DestinationDiskExists(path=instance_dir)
39            os.mkdir(instance_dir)
40
41        if not is_shared_block_storage:
42            # Ensure images and backing files are present.
43            self._create_images_and_backing(context, instance,
44                                            instance_dir, disk_info)
45
46        if not (is_block_migration or is_shared_instance_path):
47            # NOTE(angdraug): when block storage is shared between source and
48            # destination and instance path isn't (e.g. volume backed or rbd
49            # backed instance), instance path on destination has to be prepared
50

```

```

51         # Touch the console.log file, required by libvirt.
52         console_file = self._get_console_log_path(instance)
53         libvirt_utils.file_open(console_file, 'a').close()
54
55         # if image has kernel and ramdisk, just download
56         # following normal way.
57         self._fetch_instance_kernel_ramdisk(context, instance)
58
59         # Establishing connection to volume server.
60         block_device_mapping = driver.block_device_info_get_mapping(
61             block_device_info)
62         for vol in block_device_mapping:
63             connection_info = vol['connection_info']
64             disk_info = blockinfo.get_info_from_bdm(
65                 CONF.libvirt.virt_type, vol)
66             self._connect_volume(connection_info, disk_info)
67
68         if is_block_migration and len(block_device_mapping):
69             # NOTE(stpierre): if this instance has mapped volumes,
70             # we can't do a block migration, since that will
71             # result in volumes being copied from themselves to
72             # themselves, which is a recipe for disaster.
73             LOG.error(
74                 _LE('Cannot block migrate instance %s with mapped volumes') %
75                 instance.uuid)
76             raise exception.MigrationError(
77                 _('Cannot block migrate instance %s with mapped volumes') %
78                 instance.uuid)
79
80         # We call plug_vifs before the compute manager calls
81         # ensure_filtering_rules_for_instance, to ensure bridge is set up
82         # Retry operation is necessary because continuously request comes,
83         # concurrent request occurs to iptables, then it complains.
84         max_retry = CONF.live_migration_retry_count
85         for cnt in range(max_retry):
86             try:
87                 self.plug_vifs(instance, network_info)
88                 break
89             except processutils.ProcessExecutionError:
90                 if cnt == max_retry - 1:
91                     raise
92                 else:
93                     LOG.warn(_LW('plug_vifs() failed %(cnt)d. Retry up to '
94                                 '%(max_retry)d. '),
95                             {'cnt': cnt,
96                              'max_retry': max_retry},
97                             instance=instance)
98                     greenthread.sleep(1)
99
100         res_data = {'graphics_listen_addrs': {}}
101         res_data['graphics_listen_addrs']['vnc'] = CONF.vncserver_listen
102         res_data['graphics_listen_addrs']['spice'] = CONF.spice.server_listen
103
104         return res_data

```

5.2.2 源主机上的操作: live_migration

```

1  # nova/compute/manager.py ComputeManager.live_migration()
2  def live_migration(self, context, dest, instance, block_migration,
3                    migrate_data):
4      """Executing live migration.
5
6      :param context: security context
7      :param instance: a nova.objects.instance.Instance object
8      :param dest: destination host
9      :param block_migration: if true, prepare for block migration
10     :param migrate_data: implementation specific params
11
12     """
13
14     # NOTE(danms): since instance is not the first parameter, we can't
15     # use @object_compat on this method. Since this is the only example,
16     # we do this manually instead of complicating the decorator
17     if not isinstance(instance, obj_base.NovaObject):
18         expected = ['metadata', 'system_metadata',
19                    'security_groups', 'info_cache']
20         instance = objects.Instance._from_db_object(
21             context, objects.Instance(), instance,
22             expected_attrs=expected)
23
24     # Create a local copy since we'll be modifying the dictionary
25     migrate_data = dict(migrate_data or {})
26     try:
27         if block_migration:
28             block_device_info = self._get_instance_block_device_info(
29                 context, instance)
30             disk = self.driver.get_instance_disk_info(
31                 instance.name, block_device_info=block_device_info)
32         else:
33             disk = None
34
35         pre_migration_data = self.compute_rpcapi.pre_live_migration(
36             context, instance,
37             block_migration, disk, dest, migrate_data)
38         migrate_data['pre_live_migration_result'] = pre_migration_data
39
40     except Exception:
41         with excutils.save_and_reraise_exception():
42             LOG.exception(_LE('Pre live migration failed at %s'),
43                           dest, instance=instance)
44             self._rollback_live_migration(context, instance, dest,
45                                           block_migration, migrate_data)
46
47     # Executing live migration
48     # live_migration might raises exceptions, but
49     # nothing must be recovered in this version.
50     self.driver.live_migration(context, instance, dest,
51                               self._post_live_migration,
52                               self._rollback_live_migration,
53                               block_migration, migrate_data)

```

```

1  # nova/virt/libvirt/driver.py LibvirtDriver.live_migration()
2  def live_migration(self, context, instance, dest,
3                    post_method, recover_method, block_migration=False,
4                    migrate_data=None):
5      """Spawning live_migration operation for distributing high-load.
6
7      :param context: security context
8      :param instance:
9          nova.db.sqlalchemy.models.Instance object
10         instance object that is migrated.
11      :param dest: destination host
12      :param post_method:
13         post operation method.
14         expected nova.compute.manager._post_live_migration.
15      :param recover_method:
16         recovery method when any exception occurs.
17         expected nova.compute.manager._rollback_live_migration.
18      :param block_migration: if true, do block migration.
19      :param migrate_data: implementation specific params
20
21      """
22
23      # 'dest' will be substituted into 'migration_uri' so ensure
24      # it does't contain any characters that could be used to
25      # exploit the URI accepted by libvirt
26      if not libvirt_utils.is_valid_hostname(dest):
27          raise exception.InvalidHostname(hostname=dest)
28
29      greenthread.spawn(self._live_migration, context, instance, dest,
30                       post_method, recover_method, block_migration,
31                       migrate_data)

```

```

1  # nova/virt/libvirt/driver.py LibvirtDriver._live_migration()
2  def _live_migration(self, context, instance, dest, post_method,
3                    recover_method, block_migration=False,
4                    migrate_data=None):
5      """Do live migration.
6
7      :param context: security context
8      :param instance:
9          nova.db.sqlalchemy.models.Instance object
10         instance object that is migrated.
11      :param dest: destination host
12      :param post_method:
13         post operation method.
14         expected nova.compute.manager._post_live_migration.
15      :param recover_method:
16         recovery method when any exception occurs.
17         expected nova.compute.manager._rollback_live_migration.
18      :param block_migration: if true, do block migration.
19      :param migrate_data: implementation specific params
20
21      """
22
23      # Do live migration.
24      try:

```

```

24         if block_migration:
25             flaglist = CONF.libvirt.block_migration_flag.split(',')
26         else:
27             flaglist = CONF.libvirt.live_migration_flag.split(',')
28         flagvals = [getattr(libvirt, x.strip()) for x in flaglist]
29         logical_sum = reduce(lambda x, y: x | y, flagvals)
30
31         dom = self._lookup_by_name(instance["name"])
32
33         pre_live_migrate_data = (migrate_data or {}).get(
34             'pre_live_migration_result', {})
35         listen_addrs = pre_live_migrate_data.get('graphics_listen_addrs')
36
37         migratable_flag = getattr(libvirt, 'VIR_DOMAIN_XML_MIGRATABLE',
38                                   None)
39
40         if migratable_flag is None or listen_addrs is None:
41             self._check_graphics_addresses_can_live_migrate(listen_addrs)
42             dom.migrateToURI(CONF.libvirt.live_migration_uri % dest,
43                             logical_sum,
44                             None,
45                             CONF.libvirt.live_migration_bandwidth)
46         else:
47             old_xml_str = dom.XMLDesc(migratable_flag)
48             new_xml_str = self._correct_listen_addr(old_xml_str,
49                                                     listen_addrs)
50
51             try:
52                 dom.migrateToURI2(CONF.libvirt.live_migration_uri % dest,
53                                   None,
54                                   new_xml_str,
55                                   logical_sum,
56                                   None,
57                                   CONF.libvirt.live_migration_bandwidth)
58             except libvirt.libvirtError as ex:
59                 # NOTE(mriedem): There is a bug in older versions of
60                 # libvirt where the VIR_DOMAIN_XML_MIGRATABLE flag causes
61                 # virDomainDefCheckABIStability to not compare the source
62                 # and target domain xml's correctly for the CPU model.
63                 # We try to handle that error here and attempt the legacy
64                 # migrateToURI path, which could fail if the console
65                 # addresses are not correct, but in that case we have the
66                 # _check_graphics_addresses_can_live_migrate check in place
67                 # to catch it.
68                 # TODO(mriedem): Remove this workaround when
69                 # Red Hat BZ #1141838 is closed.
70                 error_code = ex.get_error_code()
71                 if error_code == libvirt.VIR_ERR_CONFIG_UNSUPPORTED:
72                     LOG.warn(_LW('An error occurred trying to live '
73                                   'migrate. Falling back to legacy live '
74                                   'migrate flow. Error: %s'), ex,
75                               instance=instance)
76                 self._check_graphics_addresses_can_live_migrate(
77                     listen_addrs)
78                 dom.migrateToURI(
79                     CONF.libvirt.live_migration_uri % dest,
80                     logical_sum,

```

```
80         None,
81         CONF.libvirt.live_migration_bandwidth)
82     else:
83         raise
84
85     except Exception as e:
86         with excutils.save_and_reraise_exception():
87             LOG.error(_LE("Live Migration failure: %s"), e,
88                       instance=instance)
89             recover_method(context, instance, dest, block_migration)
90
91     post_method(context, instance, dest, block_migration,
92                migrate_data)
```