

目 录

1 snapshot 的流程

1.1 数据结构

1.1.1 req

```
1 # nova/api/openstack/wsgi.py Request
```

1.1.2 body

```
1 {u'createImage': {u'name': u'snap1', u'metadata': {}}}
```

1.2 nova-api 部分

```
1 # nova/api/openstack/compute/servers.py Controller._action_create_image()
2 def _action_create_image(self, req, id, body):
3     # id是instance id
4     """Snapshot a server instance."""
5     context = req.environ['nova.context']
6     entity = body.get("createImage", {})
7
8     image_name = entity.get("name")
9
10    ...
11
12    props = {}
13    metadata = entity.get('metadata', {})
14    common.check_img_metadata_properties_quota(context, metadata)
15    try:
16        props.update(metadata)
17    except ValueError:
18        msg = _("Invalid metadata")
19        raise exc.HTTPBadRequest(explanation=msg)
20
21    instance = self._get_server(context, req, id)
22
23    bdms = objects.BlockDeviceMappingList.get_by_instance_uuid(
24        context, instance.uuid)
25
26    try:
27        if self.compute_api.is_volume_backed_instance(context, instance,
28                                                       bdms):
29            img = instance['image_ref']
30            if not img:
31                properties = bdms.root_metadata(
32                    context, self.compute_api.image_api,
33                    self.compute_api.volume_api)
```

```

34         image_meta = {'properties': properties}
35     else:
36         image_meta = self.compute_api.image_api.get(context, img)
37
38         image = self.compute_api.snapshot_volume_backed(
39             context,
40             instance,
41             image_meta,
42             image_name,
43             extra_properties=props)
44     else:
45         image = self.compute_api.snapshot(context,
46             instance,
47             image_name,
48             extra_properties=props)
49 except exception.InstanceInvalidState as state_error:
50     common.raise_http_conflict_for_instance_invalid_state(state_error,
51         'createImage')
52 except exception.Invalid as err:
53     raise exc.HTTPBadRequest(explanation=err.format_message())
54
55 # build location of newly-created image entity
56 image_id = str(image['id'])
57 url_prefix = self._view_builder._update_glance_link_prefix(
58     req.application_url)
59 image_ref = os.path.join(url_prefix,
60     context.project_id,
61     'images',
62     image_id)
63
64 resp = webob.Response(status_int=202)
65 resp.headers['Location'] = image_ref
66 return resp

```

```

1 # nova/compute/api.py API.snapshot()
2 def snapshot(self, context, instance, name, extra_properties=None):
3     """Snapshot the given instance.
4
5     :param instance: nova.db.sqlalchemy.models.Instance
6     :param name: name of the snapshot
7     :param extra_properties: dict of extra image properties to include
8                             when creating the image.
9     :returns: A dict containing image metadata
10    """
11    image_meta = self._create_image(context, instance, name,
12        'snapshot',
13        extra_properties=extra_properties)
14
15    # NOTE(comstud): Any changes to this method should also be made
16    # to the snapshot_instance() method in nova/cells/messaging.py
17    instance.task_state = task_states.IMAGE_SNAPSHOT_PENDING
18    instance.save(expected_task_state=[None])
19
20    self.compute_rpcapi.snapshot_instance(context, instance,
21        image_meta['id'])

```

```

22
23     return image_meta

```

```

1  # nova/compute/rpcapi.py ComputeAPI.snapshot_instance()
2  def snapshot_instance(self, ctxt, instance, image_id):
3      version = '3.0'
4      cctxt = self.client.prepare(server=_compute_host(None, instance),
5                                  version=version)
6      cctxt.cast(ctxt, 'snapshot_instance',
7                  instance=instance,
8                  image_id=image_id)

```

1.3 nova-compute 部分

```

1  # nova/compute/manager.py ComputeManager._snapshot_instance()
2  def snapshot_instance(self, context, image_id, instance):
3      """Snapshot an instance on this host.
4
5      :param context: security context
6      :param instance: a nova.objects.instance.Instance object
7      :param image_id: glance.db.sqlalchemy.models.Image.Id
8      """
9      # NOTE(dave-mcnally) the task state will already be set by the api
10     # but if the compute manager has crashed/been restarted prior to the
11     # request getting here the task state may have been cleared so we set
12     # it again and things continue normally
13     try:
14         instance.task_state = task_states.IMAGE_SNAPSHOT
15         instance.save(
16             expected_task_state=task_states.IMAGE_SNAPSHOT_PENDING)
17     except exception.InstanceNotFound:
18         # possibility instance no longer exists, no point in continuing
19         LOG.debug("Instance not found, could not set state %s "
20                 "for instance.",
21                 task_states.IMAGE_SNAPSHOT, instance=instance)
22         return
23
24     except exception.UnexpectedDeletingTaskStateError:
25         LOG.debug("Instance being deleted, snapshot cannot continue",
26                 instance=instance)
27         return
28
29     self._snapshot_instance(context, image_id, instance,
30                             task_states.IMAGE_SNAPSHOT)

```

```

1  def _snapshot_instance(self, context, image_id, instance,
2                          expected_task_state):
3      context = context.elevated()
4
5      current_power_state = self._get_power_state(context, instance)
6      try:

```

```

7         instance.power_state = current_power_state
8         instance.save()
9
10        LOG.audit(_('instance snapshotting'), context=context,
11                  instance=instance)
12
13        if instance.power_state != power_state.RUNNING:
14            state = instance.power_state
15            running = power_state.RUNNING
16            LOG.warn(_('trying to snapshot a non-running instance: '
17                      '(state: %(state)s expected: %(running)s)'),
18                    {'state': state, 'running': running},
19                    instance=instance)
20
21        self._notify_about_instance_usage(
22            context, instance, "snapshot.start")
23
24        def update_task_state(task_state,
25                              expected_state=expected_task_state):
26            instance.task_state = task_state
27            instance.save(expected_task_state=expected_state)
28
29        self.driver.snapshot(context, instance, image_id,
30                             update_task_state)
31
32        instance.task_state = None
33        instance.save(expected_task_state=task_states.IMAGE_UPLOADING)
34
35        self._notify_about_instance_usage(context, instance,
36                                          "snapshot.end")
37    except (exception.InstanceNotFound,
38            exception.UnexpectedDeletingTaskStateError):
39        # the instance got deleted during the snapshot
40        # Quickly bail out of here
41        msg = 'Instance disappeared during snapshot'
42        LOG.debug(msg, instance=instance)
43        try:
44            image_service = glance.get_default_image_service()
45            image = image_service.show(context, image_id)
46            if image['status'] != 'active':
47                image_service.delete(context, image_id)
48        except Exception:
49            LOG.warning(_("Error while trying to clean up image %s"),
50                       image_id, instance=instance)
51    except exception.ImageNotFound:
52        instance.task_state = None
53        instance.save()
54        msg = _("Image not found during snapshot")
55        LOG.warn(msg, instance=instance)

```

```

1    def snapshot(self, context, instance, image_id, update_task_state):
2        """Create snapshot from a running VM instance.
3
4        This command only works with qemu 0.14+
5        """

```

```

6         try:
7             virt_dom = self._lookup_by_name(instance['name'])
8         except exception.InstanceNotFound:
9             raise exception.InstanceNotRunning(instance_id=instance['uuid'])
10
11         base_image_ref = instance['image_ref']
12
13         base = compute_utils.get_image_metadata(
14             context, self._image_api, base_image_ref, instance)
15
16         snapshot = self._image_api.get(context, image_id)
17
18         disk_path = libvirt_utils.find_disk(virt_dom)
19         source_format = libvirt_utils.get_disk_type(disk_path)
20
21         image_format = CONF.libvirt.snapshot_image_format or source_format
22
23         # NOTE(bfilippov): save lvm and rbd as raw
24         if image_format == 'lvm' or image_format == 'rbd':
25             image_format = 'raw'
26
27         metadata = self._create_snapshot_metadata(base,
28                                                    instance,
29                                                    image_format,
30                                                    snapshot['name'])
31
32         snapshot_name = uuid.uuid4().hex
33
34         state = LIBVIRT_POWER_STATE[virt_dom.info()[0]]
35
36         # NOTE(rmk): Live snapshots require QEMU 1.3 and Libvirt 1.0.0.
37         # These restrictions can be relaxed as other configurations
38         # can be validated.
39         # NOTE(dgenin): Instances with LVM encrypted ephemeral storage require
40         # cold snapshots. Currently, checking for encryption is
41         # redundant because LVM supports only cold snapshots.
42         # It is necessary in case this situation changes in the
43         # future.
44         if (self._has_min_version(MIN_LIBVIRT_LIVESNAPSHOT_VERSION,
45                                   MIN_QEMU_LIVESNAPSHOT_VERSION,
46                                   REQ_HYPERVISOR_LIVESNAPSHOT)
47             and source_format not in ('lvm', 'rbd')
48             and not CONF.ephemeral_storage_encryption.enabled):
49             live_snapshot = True
50             # Abort is an idempotent operation, so make sure any block
51             # jobs which may have failed are ended. This operation also
52             # confirms the running instance, as opposed to the system as a
53             # whole, has a new enough version of the hypervisor (bug 1193146).
54             try:
55                 virt_dom.blockJobAbort(disk_path, 0)
56             except libvirt.libvirtError as ex:
57                 error_code = ex.get_error_code()
58                 if error_code == libvirt.VIR_ERR_CONFIG_UNSUPPORTED:
59                     live_snapshot = False
60                 else:
61                     pass

```

```

62         else:
63             live_snapshot = False
64
65         # NOTE(rmk): We cannot perform live snapshots when a managedSave
66         #             file is present, so we will use the cold/legacy method
67         #             for instances which are shutdown.
68         if state == power_state.SHUTDOWN:
69             live_snapshot = False
70
71         # NOTE(dkang): managedSave does not work for LXC
72         if CONF.libvirt.virt_type != 'lxc' and not live_snapshot:
73             if state == power_state.RUNNING or state == power_state.PAUSED:
74                 self._detach_pci_devices(virt_dom,
75                                           pci_manager.get_instance_pci_devs(instance))
76                 self._detach_sriov_ports(context, instance, virt_dom)
77                 virt_dom.managedSave(0)
78
79         snapshot_backend = self.image_backend.snapshot(instance,
80                                                         disk_path,
81                                                         image_type=source_format)
82
83         if live_snapshot:
84             LOG.info(_LI("Beginning live snapshot process"),
85                     instance=instance)
86         else:
87             LOG.info(_LI("Beginning cold snapshot process"),
88                     instance=instance)
89
90         update_task_state(task_state=task_states.IMAGE_PENDING_UPLOAD)
91         snapshot_directory = CONF.libvirt.snapshots_directory
92         fileutils.ensure_tree(snapshot_directory)
93         with utils.tmpdir(dir=snapshot_directory) as tmpdir:
94             try:
95                 out_path = os.path.join(tmpdir, snapshot_name)
96                 if live_snapshot:
97                     # NOTE(xqueralto): libvirt needs o+x in the temp directory
98                     os.chmod(tmpdir, 0o701)
99                     self._live_snapshot(virt_dom, disk_path, out_path,
100                                         image_format)
101                 else:
102                     snapshot_backend.snapshot_extract(out_path, image_format)
103             finally:
104                 new_dom = None
105                 # NOTE(dkang): because previous managedSave is not called
106                 #             for LXC, _create_domain must not be called.
107                 if CONF.libvirt.virt_type != 'lxc' and not live_snapshot:
108                     if state == power_state.RUNNING:
109                         new_dom = self._create_domain(domain=virt_dom)
110                     elif state == power_state.PAUSED:
111                         new_dom = self._create_domain(domain=virt_dom,
112                                                         launch_flags=libvirt.VIR_DOMAIN_START_PAUSED)
113                     if new_dom is not None:
114                         self._attach_pci_devices(new_dom,
115                                                   pci_manager.get_instance_pci_devs(instance))
116                         self._attach_sriov_ports(context, instance, new_dom)
117                 LOG.info(_LI("Snapshot extracted, beginning image upload"),

```

```
118         instance=instance)
119
120     # Upload that image to the image service
121
122     update_task_state(task_state=task_states.IMAGE_UPLOADING,
123                       expected_state=task_states.IMAGE_PENDING_UPLOAD)
124     with libvirt_utils.file_open(out_path) as image_file:
125         self._image_api.update(context,
126                                image_id,
127                                metadata,
128                                image_file)
129     LOG.info(_LI("Snapshot image upload complete"),
130             instance=instance)
```